

*European Conference on Computer Systems '20 (EuroSys)*

# A Fault-Tolerance Shim for Serverless Computing

**2021.05.11**

**컴퓨터학과 김명현**

# Table of Contents

---

1. Introduction
2. Background and Motivation
3. Achieving Atomicity
4. Scaling AFT
5. Garbage Collection
6. Evaluation
7. Related Work
8. Conclusion and Future Work

---

# Section 1

# “Introduction”

# 1. Introduction

---

## Serverless Computing

: an cloud computing architecture that does **not require developers to manage servers**

- users upload code and select trigger events (e.g., *API invocation, file upload*)
- than the cloud provider **transparently** manages deployment, scaling and billing.
- users are charged for resource time *used*

# 1. Introduction

---

## Serverless Computing

### **FaaS (Function-as-a-Service)**

- FaaS platforms allow user to construct applications in high-level languages
    - ex) AWS Lambda, Google Cloud Functions, Azure Function
  - there is a requirement that programs be **stateless**
    - because requests are not guaranteed to be routed to any particular instance of a program
- ⇒ applications must **be purely functional** or **modify state in a shared storage system**

# 1. Introduction

---

## Serverless Computing

### **Fault Tolerance in FaaS**

: FaaS platform provide fault tolerance with ***retries***

- if functions fail, FaaS platforms retry it automatically
- and clients will re-issue requests after a timeout.
  
- retries ensure that functions are executed at-least once
  
- developers are recommended to write **idempotent** programs

# 1. Introduction

---

## Serverless Computing

### **Idempotence**

- idempotent functions can be applied multiple times  
without changing the result beyond the initial application  
ex)  $f(f(x)) = f(x)$
- retry-based & idempotent execution would seem to guarantee exactly once execution
- but, this idempotence requirement is unreasonable for programmers  
and also insufficient to guarantee exactly once execution

# 1. Introduction

---

## Serverless Computing

### Idempotence

- function  $f$  : writes two keys,  $k$  and  $l$ , to storage
  - if  $f$  fails between the writes of  $k$  and  $l$ ,  
parallel requests might read the new version of  $k$  while reading an old version of  $l$
- ⇒ developers are forced to explicitly reason about the correctness of read operations



# 1. Introduction

---

## Serverless Computing

### **Solution**

- a simple solution is to use serializable transactions
- functions have well-defined beginnings and endings  $\Rightarrow$  guarantee atomicity
- but strongly consistent transactional systems have **scaling bottlenecks**

# 1. Introduction

---

## Serverless Computing

### **Atomic Fault-Tolerance Shim**

- **AFT** : Atomic Fault Tolerance shim between FaaS platforms and a cloud storage engine
- during a single logical request, updates written to storage are buffered by AFT
- and atomically committed at the end of request
- AFT enforces the read atomic isolation guarantee,
- ensuring that transactions never see partial side effects and only read data from sets of atomic transactions
- AFT only relies on the storage engine for durability

# 1. Introduction

## Serverless Computing

### Atomic Fault-Tolerance Shim

- **AFT** : Atomic Fault-Tolerance

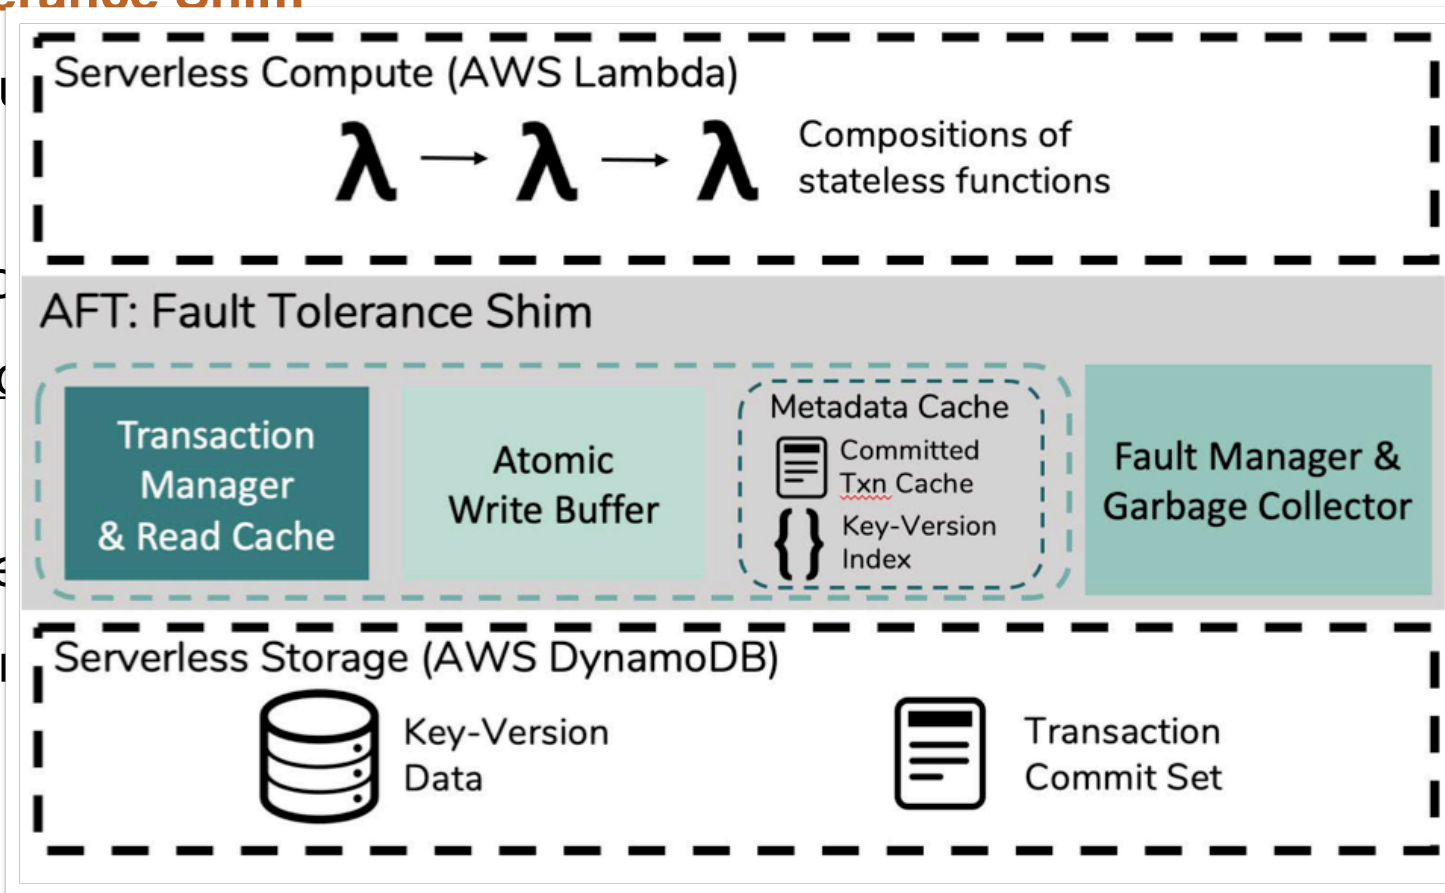
- during a single local execution

- and atomically commit

- AFT enforces the

- ensuring that transaction  
transactions

- AFT only relies on the storage engine for durability



Storage engine

Storage sets of atomic

---

## Section 2

# “Background and Motivation”

## 2. Background and Motivation

---

### Read Atomic Isolation (RA)

- RA aims to ensure that transactions do not consider partial effects of other transactions
- "A system provides RA if it prevents fractured reads anomalies and also prevents transactions from reading uncommitted, aborted, or intermediate data\*."



### **Fractured Read**

- "transaction  $T_i$  writes  $x_m$  and  $y_n$ , and  $T_j$  reads  $x_m$  and  $y_k$ . ( $k < n$ )"

\* Peter Bailis, et al. (2014)

## 2. Background and Motivation

---

### Challenges and Motivation

#### **Serverless Applications**

- serverless applications typically consist of multiple functions
- each requests are modeled as a linear composition of functions
- AFT must ensure atomic read/writes across all function in the composition,  
each function might be executed on a different machine
- and must ensure that retries guarantee idempotence

---

## Section 3

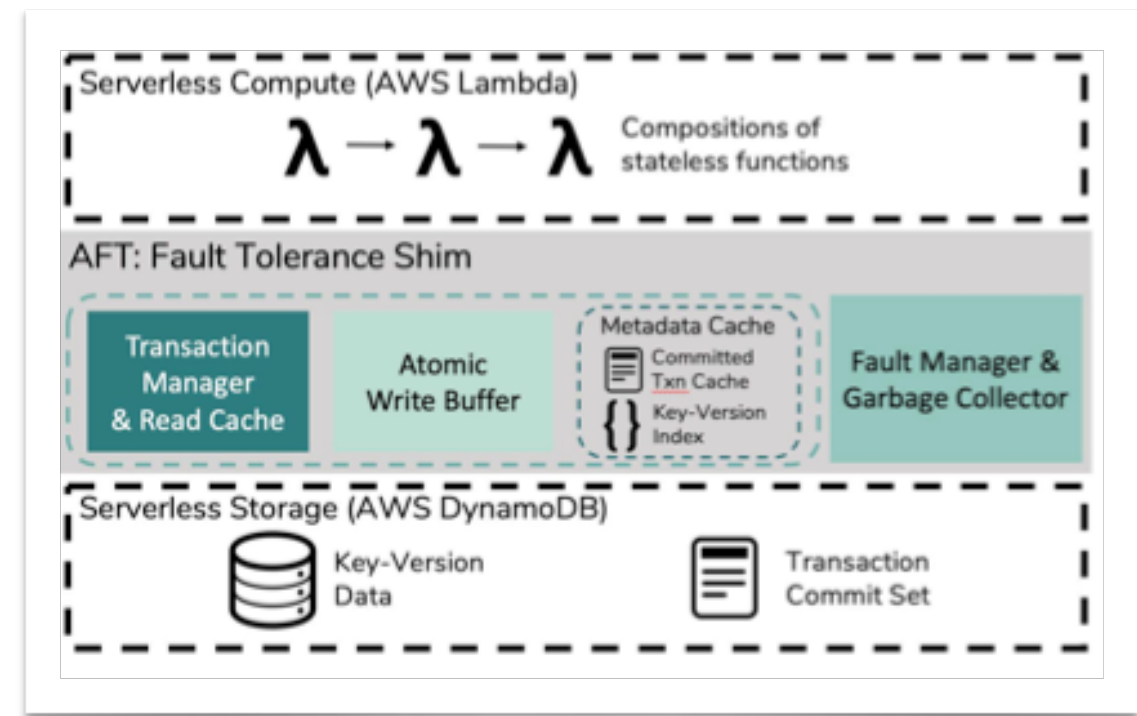
# “Achieving Atomicity”

### 3. Achieving Atomicity

#### Architecture and API

**Transaction** : each logical request

- transaction might span multiple FaaS functions





### 3. Achieving Atomicity

---

#### Architecture and API

**Transaction** : each logical request

- transaction might span multiple FaaS functions
- a new transaction begins when a client calls *StartTransaction()*
  - the transaction is assigned globally unique UUID
  - the ID of transaction (*txid*) consists of  $\langle timestamp, uuid \rangle$  pair
- AFT uses *txid* to ensure that its updates are only persisted once  
assuring idempotence in the face of retries

API	Description
<code>StartTransaction()-&gt;txid</code>	Begins a new transaction and returns a transaction ID.
<code>Get(txid, key)-&gt;value</code>	Retrieves key in the context of the transaction keyed by txid.
<code>Put(txid, key, value)</code>	Performs an update for transaction txid.
<code>AbortTransaction(txid)</code>	Aborts transaction txid and discards any updates made by it.
<code>CommitTransaction(txid)</code>	Commits transaction txid and persists its updates; only acknowledges after all data and metadata has been persisted.

### 3. Achieving Atomicity

---

#### Architecture and API

- each transaction sends all operations to a single AFT node
- within the bounds of a transaction, clients interact with AFT by calling *Get()*, *Put()*
- when a client calls *CommitTransaction()*, AFT assigns a commit timestamp, and persists all of the transaction's updates
- if a client calls *AbortTransaction()* anytime, its all updates are aborted, and the data is deleted

API	Description
<b>StartTransaction()-&gt;txid</b>	Begins a new transaction and returns a transaction ID.
<b>Get(txid, key)-&gt;value</b>	Retrieves key in the context of the transaction keyed by txid.
<b>Put(txid, key, value)</b>	Performs an update for transaction txid.
<b>AbortTransaction(txid)</b>	Aborts transaction txid and discards any updates made by it.
<b>CommitTransaction(txid)</b>	Commits transaction txid and persists its updates; only acknowledges after all data and metadata has been persisted.

### 3. Achieving Atomicity

#### Architecture and API

##### AFT node

- consists of *transaction manager*,  
*atomic write buffer*, *local metadata cache*

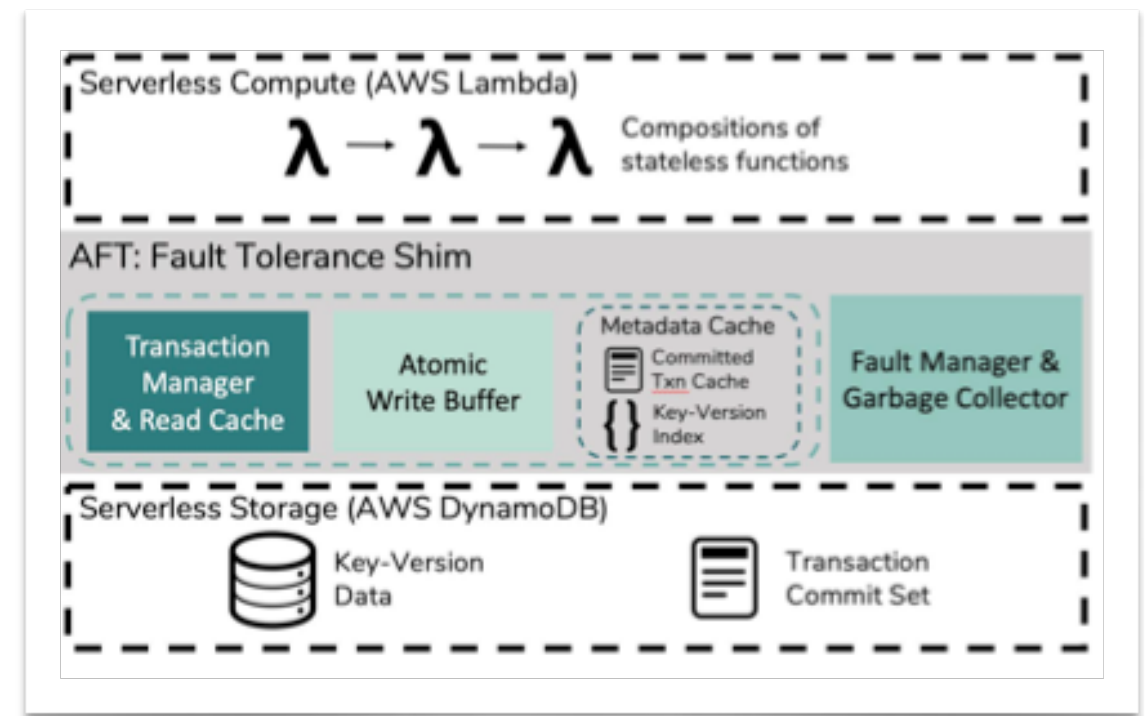
##### Transaction Manager

- tracks key versions each transaction has read,  
and enforces the read atomicity guarantee

##### Atomic Write Buffer

- gathers each transaction's writes, and persists them atomically at commit time

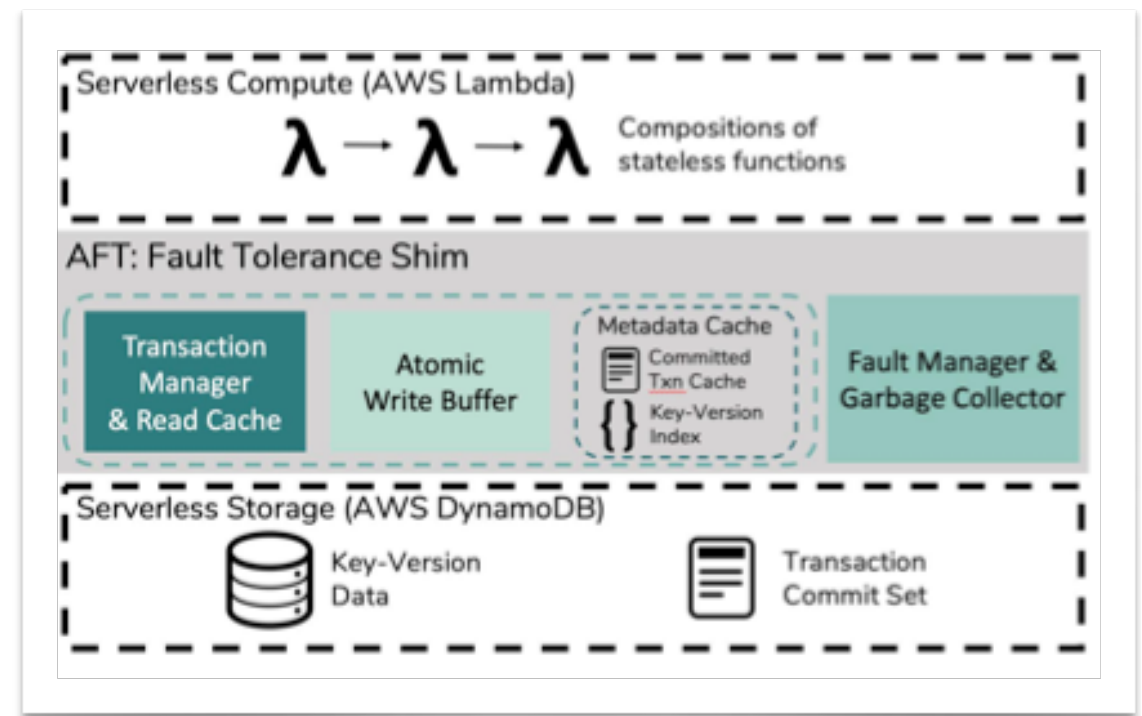
AFT maintains a *Transaction Commit Set Storage*



### 3. Achieving Atomicity

#### Architecture and API

- AFT caches recently committed transactions' ID, and locally maintains indexes of them
- also, AFT has a *data cache*, which stores values for the key versions in metadata cache
- *data cache* improves performance by avoiding storage lookups for frequently accessed versions



### 3. Achieving Atomicity

---

#### Definitions

$i$  : transaction  $T_i$  's ID (  $\langle timestamp, uuid \rangle$  )

- $T_i$  is newer than  $T_j$  ( $T_i > T_j$ ) if  $i > j$

$k_i$  : version of key  $k$  written by transaction  $T_i$

$k$  : any version of that key  $k$

- key versions are hidden from users
- clients just requests reads and writes of keys
- AFT automatically determines which versions are compatible with each request

$T_i.writeset$  : set of key versions written by transaction  $T_i$

$k_i.cowritten$  : cowritten set of key versions (  $T_i.writeset == k_i.cowritten$  )

### 3. Achieving Atomicity

---

#### Definitions

Read Atomic Isolation requires preventing *dirty reads* and *fractured reads*

#### Preventing Dirty Read

- if transaction  $T_i$  reads key version  $k_j$ , transaction  $T_j$  must have successfully committed

#### Preventing Fractured Read

- each transaction's read set must form an *Atomic Readset*

### 3. Achieving Atomicity

---

#### Definitions

##### **Definition 1. Atomic Readset**

- let  $R$  be a set of key versions
- $R$  is an *Atomic Readset* if  $\forall k_i \in R, \forall l_i \in k_i.\text{cowritten}, l_j \in R \Rightarrow j \geq i$

Example)

- there are two committed transactions in storage:  
 $T_1 : \{l_1\}$  and  $T_2 : \{k_2, l_2\}$
- if a new transaction  $T_n$  first requests  $k$  and reads  $k_2$ ,  
a subsequent read of  $l$  must return  $l_2$  or some newer version of  $l$

### 3. Achieving Atomicity

---

#### Definitions

##### **Read-Your-Writes**

- transaction must read the most recent version of a key it previously wrote

Example)

- if a transaction  $T_i$  wrote key version  $k_{i_1}$ , following read should return  $k_{i_1}$
- then if  $T_i$  writes  $k_{i_2}$ , future reads will return  $k_{i_2}$

##### **Repeatable Read**

- transaction should view the same key version if it requests the same key repeatedly

Example)

- if  $T_i$  reads version  $k_j$  then later requests a read of  $k$  again, it should read  $k_j$  unless it wrote  $k_i$



### 3. Achieving Atomicity

---

#### Preventing Dirty Reads

##### **AFT's write-ordering protocol**

- Atomic Write Buffer sequesters all updates for each transaction
- when *CommitTransaction()* is called, AFT first writes transaction's updates to storage
- after all updates have successfully been persisted,  
AFT writes the transaction's write set, timestamp, and UUID to the Transaction Commit Set
- only after the Commit Set updated, AFT make the transaction's data visible to other requests
- if *AbortTransaction()* is called, updates are simply deleted from Atomic Write Buffer

### 3. Achieving Atomicity

---

#### Preventing Dirty Reads

##### **AFT's write-ordering protocol**

- AFT doesn't overwrite keys in place (each key version is mapped to a **unique** storage key)
  - ⇒ this increases AFT's storage and metadata footprints

### 3. Achieving Atomicity

---

#### Preventing Dirty Reads - Atomic Fault Tolerance

##### **AFT's write-ordering protocol**

- when an application function fails, none of its updates will be persisted
- and the transaction will be aborted after a timeout
  
- if the function retries, it can use the same transaction ID to continue the transaction

### 3. Achieving Atomicity

---

#### Preventing Fractured Reads

##### **AFT's atomic read protocol**

- after every consecutive read, the set of key versions read forms an Atomic Readset (Definition 1)
- algorithm 1 ensures that reads are only issued from already-committed transaction  
by using the local cache of committed transaction metadata and recent key versions

### 3. Achieving Atomicity

#### Preventing Fractured Reads

#### AFT's atomic read protocol - Algorithm 1

line 3-5

- the lower bound of *target* is computed by selecting the largest transaction ID in *R* that modified *k*

**Algorithm 1** AtomicRead: For a key *k*, return a key version *k<sub>j</sub>* such that the read set *R* combined with *k<sub>j</sub>* does not violate Definition 1.

**Input:** *k, R, WriteBuffer, storage, KeyVersionIndex*

```
1: lower := 0 // Transaction ID lower bound.
2: // Lines 3-5 check case (1) of the inductive proof.
3: for li ∈ R do
4:   if k ∈ li.cowritten then // We must read kj such that j ≥ i.
5:     lower = max(lower, i)
6: // Get the latest version of k that we are aware of.
7: latest := max(KeyVersionIndex[k])
8: if latest == None ∧ lower == 0 then
9:   return NULL
10: target := None // The version of k we want to read.
11: candidateVersions := sort(filter(KeyVersionIndex[k], kv.tid ≥ lower)) // Get all versions of k at least as new as lower.
12: // Loop through versions of k in reverse timestamp order — lines 13-23 check case (2) of the inductive proof.
13: for t ∈ candidateVersions.reverse() do
14:   valid := True
15:   for li ∈ kt.cowritten do
16:     if lj ∈ R ∧ j < t then
17:       valid := False
18:       break
19:   if valid then
20:     target = t // The most recent valid version.
21:     break
22: if target == None then
23:   return NULL
24: Rnew := R ∪ {ktarget}
25: return storage.get(ktarget), Rnew
```

### 3. Achieving Atomicity

#### Preventing Fractured Reads

#### AFT's atomic read protocol - Algorithm 1

line 13-23

- iterates all cowritten keys of candidate versions
- if any cowritten key's version is in  $R$ ,  
AFT declares the candidate version valid  
if only if the cowritten key's version is not newer  
than the version in  $R$

**Algorithm 1** AtomicRead: For a key  $k$ , return a key version  $k_j$  such that the read set  $R$  combined with  $k_j$  does not violate Definition 1.

**Input:**  $k, R, WriteBuffer, storage, KeyVersionIndex$

```
1:  $lower := 0$  // Transaction ID lower bound.
2: // Lines 3-5 check case (1) of the inductive proof.
3: for  $l_i \in R$  do
4:   if  $k \in l_i.cowritten$  then // We must read  $k_j$  such that  $j \geq i$ .
5:      $lower = \max(lower, i)$ 
6: // Get the latest version of  $k$  that we are aware of.
7:  $latest := \max(KeyVersionIndex[k])$ 
8: if  $latest == None \wedge lower == 0$  then
9:   return NULL
10:  $target := None$  // The version of  $k$  we want to read.
11:  $candidateVersions := \text{sort}(\text{filter}(KeyVersionIndex[k], kv.tid \geq lower))$  // Get all versions of  $k$  at least as new as  $lower$ .
12: // Loop through versions of  $k$  in reverse timestamp order — lines 13-23 check case (2) of the inductive proof.
13: for  $t \in candidateVersions.reverse()$  do
14:    $valid := True$ 
15:   for  $l_i \in k_t.cowritten$  do
16:     if  $l_j \in R \wedge j < t$  then
17:        $valid := False$ 
18:       break
19:   if  $valid$  then
20:      $target = t$  // The most recent valid version.
21:     break
22: if  $target == None$  then
23:   return NULL
24:  $R_{new} := R \cup \{k_{target}\}$ 
25: return  $storage.get(k_{target}), R_{new}$ 
```

### 3. Achieving Atomicity

---

#### Other Guarantees

##### **Read-Your-Writes**

- when a transaction requests a key currently stored in its own write set,  
(Atomic Write Buffer does not yet assign a commit timestamp)  
AFT simply return the data immediately

---

## Section 4

# “Scaling AFT”



## 4. Scaling AFT

---

### Introduction

- each machine has a background thread:
  - periodically gathers all transactions committed recently on this node
  - broadcasts them to all other nodes
  - listens for messages from other replicas
- when the thread receives a new commit set,
  - it adds all transactions to its local Commit Set Cache and updates its key version index
- in a distributed environment, the cost of communicating this metadata can be extremely high

## 4. Scaling AFT

---

### Pruning Commit Sets

- AFT prunes the set of transactions that each node multicasts
- **locally superseded transactions** do not need to be broadcast
  - a transaction  $T_i$  is locally superseded if  $\forall k_i \in T_i . \text{writeset}, \exists k_j \mid j > i$

## 4. Scaling AFT

---

### Pruning Commit Sets

#### **Algorithm 2 : *IsTransactionSuperseded***

- check whether  $T_i$  has been superseded
- each node's background multicast protocol checks before sending it to other replicas
- the receiving node checks whether the transaction has been superseded
  - if it is true, AFT do not merge the transaction into the metadata cache
- once a transaction is superseded, the node can safely delete the transaction metadata

**Algorithm 2** IsTransactionSuperseded: Check whether transaction  $T_i$  has been superseded—if there is a newer version of every key version written by  $T_i$ .

**Input:**  $T_i, keyVersionIndex$

```
1: for  $k_i \in T_i.writeset$  do
2:    $latest := k.latest\_tid()$ 
3:   if  $latest == i$  then
4:     return False
5: return True
```

## 4. Scaling AFT

---

### Fault Tolerance

AFT considers both **safety** and **liveness**

#### **guaranteeing safety**

- AFT relies on *write-ordering protocol*
- ensures that each node does not persist dirty data and commits are correctly acknowledged

## 4. Scaling AFT

---

### Fault Tolerance

AFT considers both **safety** and **liveness**

#### **guaranteeing liveness**

- if a replica commits a transaction, and fails before broadcasting the commit, (situation  $\alpha$ )

AFT must ensure that other replicas are still made aware of the committed data

- distributed deployments of AFT have a **fault manager**
- fault manager periodically scans the **Transaction Commit Set** in storage,  
and checks for commit records that it has not received via broadcast
- in situation  $\alpha$ , fault manager simply read the transaction's commit record again

## 4. Scaling AFT

---

### Deployment and Autoscaling

#### **Deploying AFT**

- AFT is deployed using **Kubernetes**
- fault manager run in separate Docker containers and on separate machines
- fault manager detects failed nodes and configuring their replacements

---

## Section 5

# “Garbage Collection”

## 5. Garbage Collection

---

### Introduction

these data sources cause **overheads**:

- the list of all transactions committed by AFT
- the set of key versions  
( AFT does not overwrite key versions )



## 5. Garbage Collection

---

### Local Metadata Garbage Collection

- locally garbage collect transaction metadata
- background GC process periodically sweeps all committed transactions in the metadata cache
- conditions of being garbage collected:
  - if the transaction is **superseded** (Algorithm 2)
  - if no currently-executing transactions have read from the transaction's write set
- GC remove that transaction from the Commit Set Cache

## 5. Garbage Collection

---

### Global Data Garbage Collection

- **the fault manager** also serves as a global garbage collector
  - fault manager receives commit broadcasts
- global GC executes Algorithm 2 to determine which transaction is superseded
  - generates a list of transactions considered superseded
  - asks all nodes if they have locally deleted those transactions
- when the GC process receives acknowledgements from all nodes,
  - it deletes the transaction's write and commit metadata

## 5. Garbage Collection

---

### Global Data Garbage Collection - Limitation

- if a running transaction  $T_j$  has read from  $T_i$ 's write set, GC protocol will not delete  $T_i$
- but we do not know each running transaction's full read set
- example)
  - $T_1 : \{k_1\}$ ,  $T_2 : \{l_2\}$ ,  $T_3 : \{k_3, l_3\}$
  - a transaction  $T_r$  first reads  $k_1$  then requests key  $l$
  - GC process will not delete  $T_1$ , because  $T_r$  is reading  $T_1$  and is not committed yet
  - GC process might delete  $T_2$
  - if  $T_r$  attempts to read  $l$ , it will find no valid version since  $l_c$  is invalid

---

# Section 6

# “Evaluation”

## 6. Evaluation

---

### Environment

- key design goal for AFT is to run on a **variety of cloud storage backends**
- AFT is implemented over *AWS S3*\* and *AWS DynamoDB*\*\*
- AFT is run over *Redis* (deployed via AWS ElastiCache)
- AFT is implemented in Go & Python, and runs on top of Kubernetes
- simple stateless load balancer routes request to AFT nodes in a **round-robin** fashion
- all experiments were run in the us-east-1a AWS availability zone
  - AFT node ran on a *c5.2xlarge* EC2 instance with **8 vCPUs** (4 physical cores) and **16GB RAM**

\* AWS S3 (Simple Storage Service) : large scale object storage system

\*\* AWS DynamoDB : cloud-native key-value store

## 6. Evaluation

---

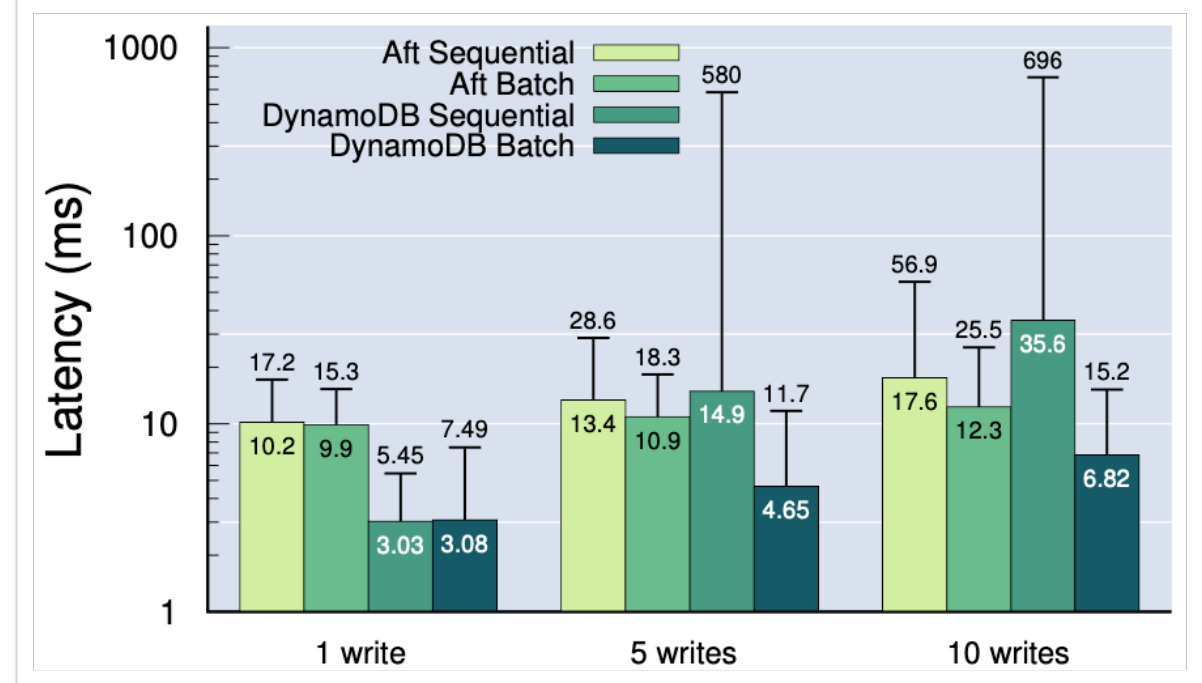
### AFT Overheads - I/O Latency

- the experiment compares the cost of writing data **directly to DynamoDB**,  
and to the cost of the same set of writes **using AFT's commit protocol**
- to isolate the overheads from FaaS platforms, we issue writes from a single thread in VM
- two baselines write directly to DynamoDB,  
one with **sequential writes** and the other with **batching**
- AFT uses batched writes by default
- but the experiment measures two configurations over AFT:  
sequential writes and single batched write

## 6. Evaluation

### AFT Overheads - I/O Latency

- the latency of DynamoDB Sequential writes increases **linearly** with the number of writes
- AFT Sequential's performance scales better than DynamoDB Sequential's
- overheads:
  - the cost of the extra network overhead by shipping data to AFT
  - the cost of writing the extra commit record



## 6. Evaluation

---

### AFT Overheads - End-to-End Latency

- the experiment compares the latency of executing transaction on **AWS Lambda with AFT**, and **without AFT** interposed between the compute and storage layers
- evaluates three storage systems: *AWS S3*, *AWS DynamoDB*, *Redis*
- each transaction is composed of 2 functions, each functions consists of 1 read and 2 writes (each reads/writes are of 4KB objects)
  - these small transactions reflects real-world CRUD applications



## 6. Evaluation

---

### AFT Overheads - End-to-End Latency

#### Consistency Anomalies

- Read Atomic Consistency Guarantee is AFT's key advantage over S3, DynamoDB and Redis
- the experiment\* measure *Read-Your-Write* anomalies (RYW) and *Fractured Read* (FR)
  - RYW : occurs when a transaction writes a key version and does not later read the same data
  - FR : occurs when a later transaction read the old data

Storage Engine	Consistency Level	RYW Anomalies	FR Anomalies
AFT	Read Atomic	0	0
S3	None	595	836
DynamoDB	None	537	779
DynamoDB	Serializable	0	115
Redis	Shard Linearizable	215	383

\* 10000 transactions

## 6. Evaluation

---

### AFT Overheads - End-to-End Latency

#### Consistency Anomalies

Storage Engine	Consistency Level	RYW Anomalies	FR Anomalies
AFT	Read Atomic	0	0
S3	None	595	836
DynamoDB	None	537	779
DynamoDB	Serializable	0	115
Redis	Shard Linearizable	215	383

#### *Redis*

- each Redis shard is linearizable but no guarantees are made across shards
- Redis shows anomalies on 6% of requests

#### *DynamoDB with transaction mode*

- this supports transaction that are either read-only or write-only
- all operations in a transaction either succeed or fail as a group
- but, there is no atomicity guarantee for transactions that span multiple functions

## 6. Evaluation


### AFT Overheads - End-to-End Latency

#### Performance

- AFT with S3 is 25% slower than plain S3
- S3 is a throughput-oriented object store
- each transaction is composed of 2 functions, which consists of 1 read and 2 writes

#### Summary


- AFT offers performance that is competitive with cloud storage engines and eliminates a significant number of anomalies



## 6. Evaluation


### Read Caching & Data Skew

- AFT's data caching, its performance effect, its interaction with skewness
- the experiment uses same workload  
2-function transaction with 2R & 1W
- *AFT-D* : DynamoDB with AFT
- *AFT-R* : Redis with AFT
- $z$  : Zipfian Distributions:
  - 1.0 (lightly contended), 1.5 (moderately contended), 2.0 (heavily contended)



## 6. Evaluation

### Scalability



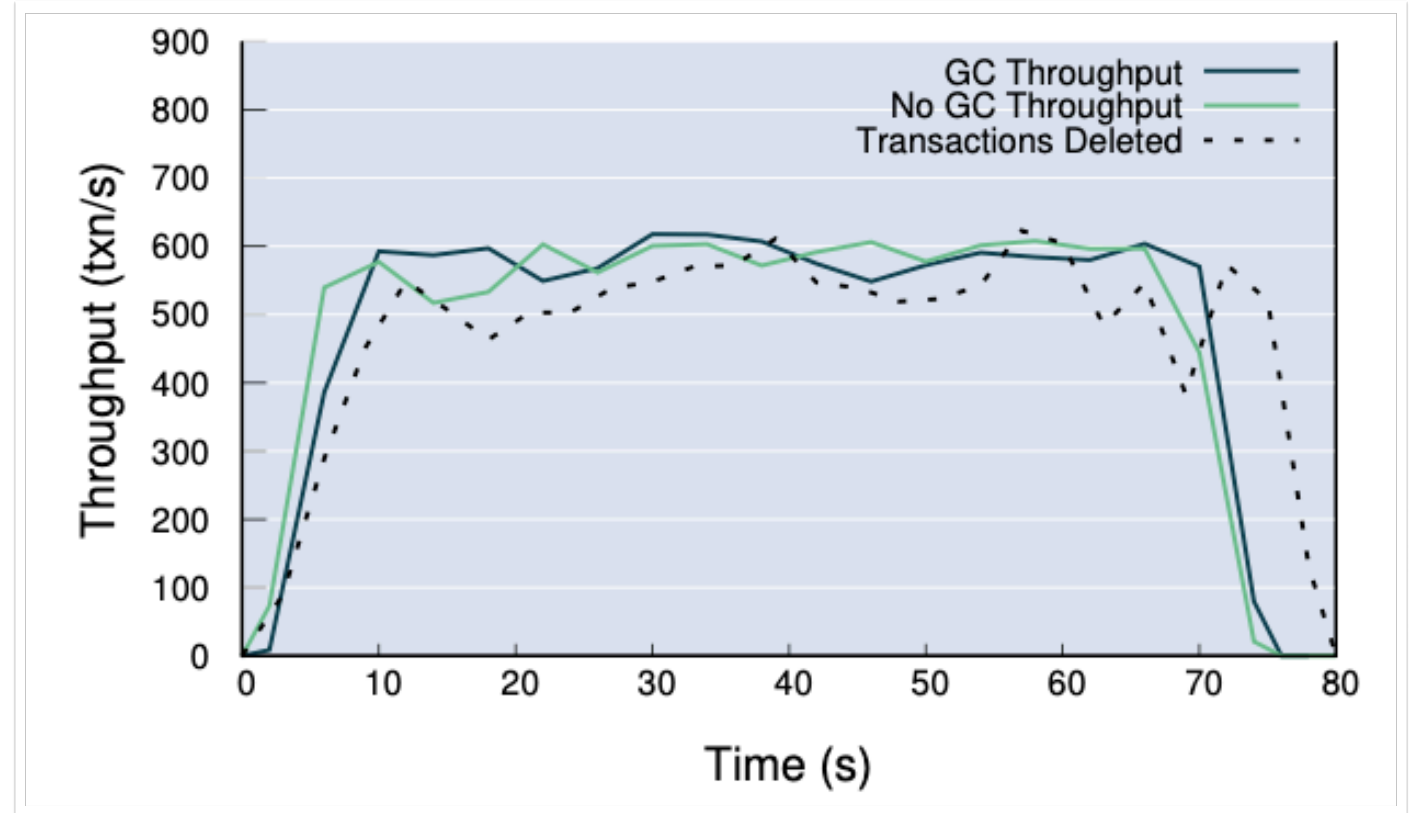
### Summary

- AFT is able to efficiency scale to thousands of transactions per second and hundreds of parallel clients within 90% of ideal throughput

## 6. Evaluation

### Garbage Collection Overheads

- run a single AFT node with 40 clients and measure throughput with GC enabled and disabled
- the cost of this GC process is that we require **separate cores** (4 AFT cores, 1 GC core)
- however, the resources allocated to GC are much smaller than the resources required to run the system




## 6. Evaluation

---

### Fault Tolerance

- this experiment measures the effect of a node failure and cost of recovering from failure  
(how long it takes for a node to cold start and join the system)
- we run AFT with 4 nodes and 200 parallel clients, and terminate a node just before 10 sec.




## 6. Evaluation

---

### Fault Tolerance

- throughput immediately drops about 16%
- the AFT management process determines that a node has failed within 5 seconds
- and it adds assigns a new node to join the cluster
- the system pre-allocate **standby nodes** to avoid having to wait for new VMs to start (starting new VM can take up to 3 minutes)






## 6. Evaluation

---

### Fault Tolerance

- over the next 45 seconds, the new node downloads the AFT docker container, updates its local metadata cache
- after 60 seconds, the node joins the cluster successfully
- the overheads from starting a new node can be mitigated by pre-downloading containers,
  - and by maintaining standby nodes with warm metadata cache



---

## Section 7

# “Related Work”

## 7. Related Work

---

### Atomic Reads

#### **Transactional Casual Consistency (TCC)\***

- guarantees atomic reads
- guarantees that reads and writes respect *Lamport's Happened-Before Relation\*\**
- but TCC model relies on fixed node at the storage layer
- also, TCC model achieve read atomicity at the storage layer

\* Syed Akbar Mehdi, *et al.* (2017)

\*\* Leslie Lamport. (1998)

## 7. Related Work

---

### Fault Tolerance

- there is a rich literature on fault-tolerance for distributed systems:
  - **checkpoint / restart** of containers or VMs\*
  - **log based replay**\*\*
- AFT is based on the simple retry model used by existing serverless platforms
- and AFT ensures atomicity and idempotence to achieve exact-once semantics

\* Wubin Li, Ali Kanso. (2015)

\*\* Stephanie Wang, *et al.* (2019)

---

## Section 8

# “Conclusion and Future Work”

## 8. Conclusion and Future Work

---

### Conclusion

#### **AFT (Atomic Fault Tolerance) Shim**

- a low-overhead fault tolerance shim for serverless computing
- AFT interposes between commodity FaaS platforms and key-value stores
- AFT achieves fault-tolerance by guaranteeing Read Atomic Isolation
- AFT guarantees exactly-once execution in the face of failures

## 8. Conclusion and Future Work

---

### Future Work

#### **Efficient Data Layout**

- AFT's key-per-version data layout works well over *DynamoDB* and *Redis*,  
but performs poorly over *S3*
- in settings with high-volume updates, the authors would like to optimize data layouts for *S3*

#### **Auto Scaling Policies**

- policies for efficient and accurate autoscaling decisions

#### **Data and Cache Partitioning**

- AFT's caching scheme is very naïve approach
- AFT will result in every node caching largely the same data, particularly for skewed workloads

---

# Q & A



---

# End of Document